An Open Deal-Level Dataset of African Agrifood Investment, 2015–2025, and the Reproducible Pipeline That Builds It

Kunal Chevuri*†, Alexander Poon*‡, Ron Mahabir§, Wisdom Mgomezulu¶,
Maction Komwa‖, Olga Gkountouna§
†High school student, Prosper High School, Prosper, TX, USA. kunalchevuri510@gmail.com
‡High school student, Edison Academy Magnet School, Edison, NJ, USA. awpoon522@gmail.com
§Geographic Data Science Lab, Department of Geography and Planning, University of Liverpool, UK
¶AGRA, Data & Analytics Unit, P.O. Box 66773, Nairobi, Kenya
‖Department of Geography and Geoinformation Science, George Mason University, Fairfax, VA, USA
*These authors contributed equally and are joint first authors.

***Abstract*—To our knowledge, there is no openly accessible, harmonized deal-level dataset that integrates development finance institution (DFI) disclosures, reported private investment rounds, and investor-level participation in African agrifood investment. DFIs post their own deals on disclosure portals, industry reports summarize venture rounds, and the rest is in press releases and trade media. Each of these records different fields, and none names every investor in a deal. We built one dataset from all three: 1,060 investor–deal rows covering 389 deals, 313 companies, 420 investors and 30 African countries from 2015 to 2025, from 121 named sources. We release it with the pipeline that builds it. The pipeline merges four portal exports, records that a large language model extracted from industry reports and we checked by hand, and investor lists we rebuilt from press coverage. It resolves entity names, assigns deterministic deal identifiers, ranks sources when they overlap, removes duplicates, and writes the same file byte for byte on every run. Every row carries its source, and every number in this paper is written by the build and checked against it by a second script. On the 73 deals that two or more sources record, the sources agree on country and year every time, on segment 94.5 percent of the time and on amount 78.0 percent of the time; 87.2 percent of rows are verified against a primary source. The main thing a user has to know is that the source decides what a record looks like. Portal deals name 1.07 investors on average, press deals 4.06, and every portal deal has a DFI in it. Comparing DFI and non-DFI deals on investor counts without accounting for this compares disclosure practice, not investment. Data and code are deposited at https://doi.org/10.5281/zenodo.23500116.**

***Index Terms*—heterogeneous data sources, data integration, entity resolution, development finance, agrifood investment**

I. Introduction

Commercial venture databases cover African agrifood thinly and oversample large rounds [1]. DFIs, the public investors that fund much of the sector, post their deals on their own portals, but a portal lists the institution's participation and rarely the other investors. AgFunder's Africa reports [2] give yearly totals and named examples, not transaction records. Company announcements and trade media name the full investor list for a round, but only for rounds someone wrote about. To study who invests with whom in this market, you have to build the record yourself.

This paper makes three contributions. The *data* contribution is an open, harmonized deal-level dataset of African agrifood investment that joins DFI disclosures, reported private rounds and investor-level participation in one record. The *methodological* contribution is the reproducible integration pipeline that builds it from those three kinds of source and writes the same file byte for byte on every run. The *measurement* contribution is evidence that the disclosure regime a record comes from systematically shapes the syndicate structure an analyst observes, which makes naive comparisons of DFI and non-DFI deals potentially misleading.

This paper describes that build. The dataset is the input to a study of DFI and private co-investment networks that is under review elsewhere [3]. That study's results are not repeated here. This paper covers the sources, the pipeline, the checks we could run on the data, and one property of the data that anyone building a similar dataset will hit: the source a record comes from decides what fields it has, and the sources are not spread evenly across the deals you want to compare.

The dataset is small in rows. The problem it solves is not volume but heterogeneity: three disclosure regimes with different fields, different coverage and different reliability, which have to be merged without losing track of which is which. The pipeline is built so that nothing in the dataset or the paper depends on a step nobody can rerun. Section VI states the rules.

II. Related data resources

Commercial venture databases cover the market unevenly and oversample large rounds [1]. Briter sells subscription data on funding deals, businesses and investors across Africa, Latin America and Asia Pacific.[1] Africa: The Big Deal maintains a curated database of African startup deals behind a newsletter.[2]

[1]https://www.briter.co/
[2]https://thebigdeal.substack.com/

Convergence holds data on past blended finance transactions, including investors, for members.[3] The OECD and the joint MDB report publish private finance mobilised by DFIs as totals by year and sector, not as deals [4], [5]. None of these gives an open, deal-level record that joins DFI portal disclosures to the venture rounds reported in the press, with the source of every row attached. That is what this dataset is.

III. Sources

Records come in through three streams (Table I).

Portal exports. We took records from four DFI disclosure portals: the IFC Project Disclosure Portal, the African Development Bank's MapAfrica, the FMO project database and the BII portfolio database.[4] None has an "agrifood in Africa" filter, so each export is a union of searches. The IFC portal was exported five times (Agribusiness and Forestry for the Africa region, two Financial Markets searches, one with an agriculture keyword, and two sets of food and agriculture sector codes) and the batches were deduplicated on project number. MapAfrica was exported for the Agriculture and Rural Development sector, all countries, full history, and cut to the window later. Cleaner scripts for the IFC and AfDB exports map them to the common schema and assign a first-pass segment from the portal's sector text. The BII portfolio export and the FMO project pages, which were scraped, were mapped by one-off notebook code and enter the build as cleaned files. Portal records are complete for the institution's own commitments and almost never name co-investors.

LLM extraction from reports. Claude (Anthropic) read the 2022, 2023 and 2024 AgFunder Africa reports and returned two tables: 131 investor–company links (an investor the report names as backing a company, with the report section) and 57 deal records (company, amount, round, year, category and country where the report gave them; 22 from the 2022 report, 29 from 2023, 6 from 2024). We checked every row against the report. After name resolution the cleaned AgFunder file holds 123 investor–company pairs; 92 are in the final dataset. Of the 31 that are not, 25 belonged to companies later found out of scope, 4 were pairs we could find no source for (the 4 known-bad pairs removed in step 1 all came from this stream), and 2 named a pitch competition rather than an investor. So 6 of the 123 links the model returned (4.9 percent) were not real investor–company relationships and 117 (95.1 percent) were; the scoping removals are not extraction errors. For 78 company names the reports did not describe well enough to classify, we researched the segment with Perplexity and then checked it against company sources; 27 of the 78 turned out to be outside scope (not African or not agrifood) and were dropped. The pipeline reads the checked classifications, not the raw model output.

Manual reconstruction. Because portals do not name syndicates, we rebuilt investor lists for named rounds by hand from company announcements, investor press releases and trade media (Disrupt Africa, TechCrunch, AgFunderNews and others), recording the URL for each row. This stream is 115 of the 121 sources and 797 of the 1,060 rows.

[3]https://www.convergence.finance/

[4]https://disclosures.ifc.org, https://mapafrica.afdb.org, https://www.fmo.nl/worldmap, https://www.bii.co.uk/en/our-impact/search-our-portfolio/; accessed 2025–2026.

TABLE I
The three streams

Stream	Sources	Rows	Deals	Investors	Inv./deal
Portal export	4	185	134	8	1.07
LLM extraction	2	78	44	20	1.39
Manual (press)	115	797	211	415	4.06
All	121	1,060	389	420	

A deal belongs to the highest-ranked stream among its rows (manual, then portal, then LLM); 57 deals have rows from more than one stream. Investors appear in more than one stream, so the column does not sum.

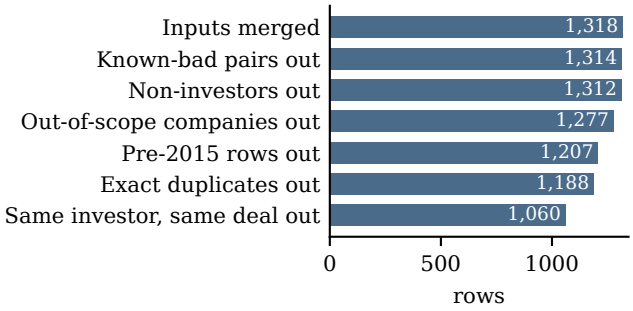


Fig. 1. Rows remaining after each step of the build.

IV. Pipeline

The build reads eight cleaned input files (1,318 rows: IFC 192, AfDB 32, FMO 7, BII 24, AgFunder 127, Disrupt Africa 647, and two manual files of 211 and 78 rows) and runs ten steps. Fig. 1 shows the row count after each step that removes rows; the build prints these counts and the figure is drawn from them.

1) **Merge and remove known-bad pairs.** Concatenate the eight inputs. Remove 4 investor–company pairs that checks during construction found no source for. The list is a pipeline step, not a one-off edit, so the removal cannot lapse when inputs are rebuilt.
2) **Entity resolution.** Map investor and company names to canonical forms with 32 investor-name rules and 16 company-name rules (for example the several spellings of "International Finance Corporation"). Set investor type with 134 rules so every row for an investor carries the same type. Remove 2 rows that name non-investors. The rules are written by hand rather than learned: with 420 investors and 313 companies, a complete rule list is small enough to read, and a similarity matcher would have needed the same hand check of every merge.
3) **Segment classification.** Pass the 142 rows still marked unknown through a keyword classifier on the company or project name (mostly IFC and AfDB project titles).

Apply the hand-checked classifications from the LLM step. Remove the 35 rows belonging to the 27 out-of-scope companies. Four rows are still unknown after this step.

4) **Country standardization.** One spelling per country. Regional and multi-country projects (65 rows) keep a regional label; 7 rows have no recoverable country.
5) **Window.** Drop the 70 rows dated before 2015. Those records are almost all single-investor portal disclosures and would change what the data measure.
6) **Deal identifiers.** Hash canonical company name, country and year. No two companies share an identifier. The key is deliberately coarse: it merges two financing events if one company raised twice in a year, and Section V measures how often that happens and what a finer key would cost. The hash gets a letter prefix because a spreadsheet once read an all-digit identifier as a number and corrupted the row.
7) **Source ranking.** When the same investor–deal pair comes from more than one stream, manual beats portal beats LLM. Manual rows have round-level detail the portals lack; portal rows are the institution's own disclosure; LLM rows were checked against a report, not a primary source. Without this rule, deduplication silently replaced verified rows with extracted ones.
8) **Deduplication.** Remove 19 exact duplicates and 128 rows that repeat an investor within a deal. One row per investor per deal remains.
9) **Type consistency.** Re-apply investor type after resolution and deduplication so no investor has two types.
10) **Canonical order.** Sort rows deterministically before writing.

The output is 1,060 rows in the 15 fields of Table II, one row per investor per deal. It covers 389 deals, 313 companies, 420 investors (27 DFIs, 393 others) and 30 countries. 202 deals name two or more investors. 330 deals (84.8 percent) have an amount. Segment takes six values (logistics 343 rows, financial services 228, production 161, agritech 134, processing 113, inputs 78; 3 unknown). Investor type takes six (private equity and venture capital 442, DFI 354, impact fund 186, blended vehicle 39, corporate 24, commercial bank 15).

TABLE II
FIELDS

Field	Content
deal_id	Hash of company, country, year; prefix "D"
company_name	Canonical company name
country	Country or regional label
investment_date	Date; 1 January where only the year is known
agrifood_segment	One of six, or unknown
funding_stage	Round or instrument where given
funding_amount_usd	Amount in USD; blank if undisclosed
investor_name	Canonical investor name
investor_type	One of six
is_lead_investor	Lead flag; set by construction, not informative
source_name	Named source
source_url	URL
extraction_method	structured_download, llm_extracted, manual
confidence_flag	1 verified at primary source; 2 likely correct
segment_justification	Reason for a hand-checked segment

V. Validation

Accuracy is established in three layers, each covering a different part of the data, and one layer is still missing.

Portal rows are the institution's own values. The 185 portal rows (17.5 percent of the dataset) were never typed by anyone. The IFC and AfDB cleaners read the raw portal exports, which are in the release, and write the cleaned inputs. Rerunning both cleaners today reproduced the cleaned files byte for byte (163 of the 185 rows; the FMO and BII inputs came from one-off notebook code and are not covered). An error in a portal row is either the institution's or a bug in code anyone can rerun.

Verification at construction. Every row carries a confidence flag. Flag 1 means the row was checked against a primary source (the disclosing institution, the company, or an investor's own release) when it was entered. Flag 2 means likely correct but not checked against a primary source. 924 rows (87.2 percent) are flag 1. By stream: 95.1 percent of portal rows, 90.6 percent of manual rows, 33.3 percent of LLM rows. The LLM rows are the least verified, which is why they rank last in step 7.

Agreement between sources. This is the one check that does not depend on our own verification. 73 deals are recorded by more than one named source. Where two or more sources give a value for the same field, Table III shows how often they match. Country and year match on every deal. Segment matches on 94.5 percent. Amount matches on 78.0 percent of the 50 deals where two sources both give one; in the 11 that differ, the larger figure is a median 3.24 times the smaller, which is what you get when one source reports total project cost and another the round size. Full dates match only 76.7 percent of the time because 55.4 percent of all rows have a year-only date stored as 1 January.

One identifier can hide two rounds. A key of company, country and year merges two financing events when a company raises twice in the same year, so we tested for that rather than assert its absence. 326 of the 389 deals carry a named round label or a day-level date and can be tested at all. Five deals, 1.3 percent of deals and 14 of the 1,060 rows, show evidence of more than one event: one with two named rounds, and four with day-level dates 58 to 194 days apart. They are listed in the release so a user can drop or split them.

We leave them merged because the fields that would separate them are themselves unevenly reported, and splitting on them costs more than it buys. Adding funding stage to the key creates 21 further deals, of which only 1 separates two named rounds; the other 20 separate a source that labelled the stage from a source that left it blank for the same round. Adding the date creates 17, almost all separating a day-dated press row from a year-dated portal row describing one deal. 1.3 percent is therefore the measured cost of a key coarse enough to survive sources that disagree about how much they say.

Re-reading rows against the source they cite. The sample

TABLE III
AGREEMENT ON DEALS RECORDED BY MORE THAN ONE SOURCE

Field	Deals compared	Match (%)
Country	73	100.0
Year	73	100.0
Full date	73	76.7
Segment	73	94.5
Amount	50	78.0

shipped with version 1.0 was drawn before the pipeline's scope removals ran, so 37 of its 100 rows are not in the released dataset at all (33 LLM-extracted, 4 manual). We redrew it from the data as released, seeded and sorted so the draw reproduces, keeping the 56 manual and 44 LLM-extracted split: the portal stream is covered by the first layer, and the LLM stream is oversampled because it is 7.4 percent of the rows and the least verified. Rates below are therefore reported per stream and never pooled.

For each row we fetched the recorded URL and tested whether the page states the company, the investor, the year and the amount the row records. Of the 56 manual rows, 40 could be re-read; 14 cite a publication rather than an article and 2 were unreachable. On those 40 the page confirms the company, the investor and the year for 39 of 40, and the amount for 26 of the 36 that carry one, the same disagreement about amounts that Table III shows between sources. The 2 rows the page did not confirm are listed in the release for a reader to judge; not finding a name in a page is weaker evidence than finding a wrong one.

None of the 44 LLM rows could be re-read. Every row in that stream cites AgFunder's research index rather than the report it was extracted from, so the release does not let a reader retrace it. That is a defect in the dataset rather than a result about the rows, and it is the first thing a later version should fix.

Still missing: an independent audit. The check above re-reads the source a row already names. It cannot catch a row pointed at the wrong article, it cannot judge a segment, and it is not independent of whoever entered the row. A person who did not build the dataset still has to find each source and read it. The redrawn sample is in the release with empty columns for exactly that, and the result will go into a later version.

VI. REPRODUCIBILITY

Five rules make the build reproducible.

Raw inputs are never edited. Corrections live in the pipeline as data files it reads (the known-bad list, the checked classifications, the name rules), so rebuilding from the raw exports reapplies every fix.

No step depends on run order or randomness. Deal identifiers are hashes, not counters. Source ranking is an explicit rule, not whichever row sorted first. Rows are sorted canonically before writing. Two runs on the same inputs give the same bytes; we checked this on a clean copy of the release (same MD5).

TABLE IV
ONE DEAL, TWO STREAMS: TRADEDEPOT (NIGERIA), 2021

Investor	Type	USD M	Date	Stream
Arcadia Funds	pe_vc	110.0	2021-12-06	manual
BII	DFI	110.0	2021-12-06	manual
IFC	DFI	110.0	2021-12-06	manual
MSA Capital	pe_vc	110.0	2021-12-06	manual
Novastar Ventures	pe_vc	110.0	2021-12-06	manual
Partech Africa	pe_vc	110.0	2021-12-06	manual
Sahel Capital	impact_fund	110.0	2021-12-06	manual
Endeavor Catalyst	pe_vc	42.0	2021-01-01	llm

Manual rows: TechCrunch, confidence flag 1. LLM row: AgFunder Africa Report 2022, flag 1. All eight rows share one deal identifier.

Every row says where it came from. Source name, URL, extraction method and confidence flag are columns, not documentation. A user can drop a stream, or a flag, and rerun.

Model output is an input, not a result. Where Claude extracted records or Perplexity researched a segment, the checked version is a file the pipeline reads, so regenerating model output changes nothing until someone re-checks it.

Every number in the paper is traced. One script writes every statistic here to CSV. A second script fails if any number in the text is missing from those CSVs. This paper passes.

VII. THE SOURCE DECIDES WHAT A RECORD LOOKS LIKE

Table IV shows one deal as two streams recorded it. TradeDepot's 2021 round reaches the dataset through 8 rows. Seven come from a TechCrunch report, dated to the day, with the round size and seven named investors, two of them DFIs. One comes from the AgFunder 2022 report, where the model found a different investor named against the company with a different amount and only a year. The build keeps both, because they are different investor–deal pairs, and the row tells you which is which. A user who wants only day-dated, verified rows can filter on two columns.

The three streams record different deals, and they record a deal differently (Fig. 2). A portal deal names 1.07 investors on average and 2.2 percent of portal deals name more than one. A press deal names 4.06 and 90.0 percent name more than one. The portal is not recording smaller syndicates. It is recording one investor: the one that runs the portal.

The streams are also not spread evenly across the comparison most users will want. Every portal deal includes a DFI, and no deal without a DFI comes through a portal. A straight comparison of DFI and non-DFI deals on any field tied to syndicate visibility therefore compares disclosure regimes, and shows DFI deals with fewer co-investors for reasons that have nothing to do with what investors did. The companion study restricts comparisons to the fields and subsets the streams share [3]. Any dataset built this way will have the same structure.

Coverage is uneven in time too (Fig. 3). Deals per year go from 15 in 2015 to 86 in 2022, then 45 in 2024 and 2025. The LLM stream puts 39 of its 44 deals in 2022, the year its reports cover in detail. Portal deals run between 7 and 20 a

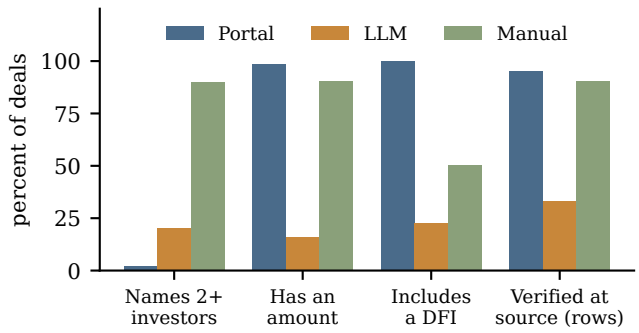


Fig. 2. Share of deals with each property, by stream.

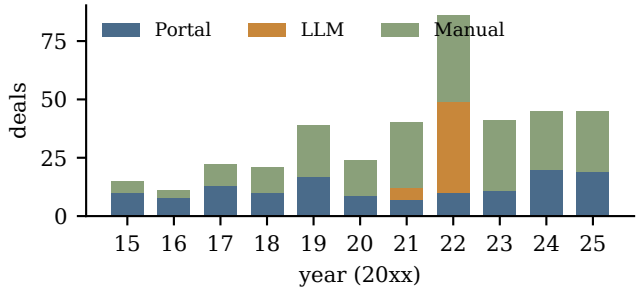


Fig. 3. Deals per year by stream.

year. Manual deals rise with press coverage. A year-by-year series from these data tracks reporting as much as investment.

VIII. Limitations and intended use

The dataset holds publicly disclosed deals. DFI investments made without announcement and private rounds nobody wrote about are missing, and the gap is one-sided: DFIs publish every commitment, private rounds appear only when reported. IFC amounts are project costs, not the IFC's commitment, so treat them as upper bounds. Over half the dates are year-only. The LLM stream is the least verified and sits mostly in one year: its deal tables concentrate there, so balancing it would need a report series with even yearly depth rather than more reports of this kind. Because `extraction_method` is a column, a user who wants a year-by-year series can drop the stream and rerun.

The data are good for asking which investors appear together, where, in what segments and when, as long as the source structure above is taken into account. On their own they are not good for comparing DFI and non-DFI deals on fields that depend on who got named, or for time series of private activity.

Data availability

The dataset is one UTF-8 CSV file of 1,060 rows and 15 columns (Table II). It is deposited on Zenodo at https://doi.org/10.5281/zenodo.23500116, a concept DOI that always resolves to the newest version, together with the eight cleaned inputs, the raw portal exports, the pipeline, the redrawn validation sample, the row-level result of the source re-check, and the scripts that produce and check every number here: data under CC BY 4.0, code under MIT. Corrections, including the result of the random audit, will go into numbered versions at the same DOI; errors can be reported to the first author. The dataset contains organisations only, no personal data. Loading it and keeping only verified, day-dated rows takes three lines:

```
df = pd.read_csv("phase1_combined.csv")
df = df[df.confidence_flag == 1]
df = df[~df.investment_date.str.endswith("-01-01")]
```

Acknowledgments

During the preparation of this work the authors used Claude (Anthropic) to assist with drafting and revising text, with literature searches, and with parts of the analysis and verification code, and used Claude and Perplexity for the extraction and classification steps described in Section III, whose output the authors checked against the sources before use. All cited sources were verified by the authors against the original publications. The authors reviewed and edited all content and take full responsibility for the content of the publication.

References

[1] M. Da Rin, T. Hellmann, and M. Puri, "A survey of venture capital research," National Bureau of Economic Research, Cambridge, MA, NBER Working Paper 17523, 2011.
[2] AgFunder, "Africa AgriFoodTech investment report 2023," AgFunder, San Francisco, CA, Tech. Rep., 2023.
[3] K. Chevuri, A. Poon, R. Mahabir, W. Mgomezulu, M. Komwa, and O. Gkountouna, "Do development finance institutions crowd in private capital? Network evidence from African agrifood investment, 2015–2025," 2026, manuscript under review.
[4] OECD, "Tracking private finance mobilisation: Latest trends and ways forward," OECD Publishing, Paris, OECD Policy Brief, 2025.
[5] Multilateral Development Banks and Development Finance Institutions, "Mobilization of private finance by multilateral development banks and development finance institutions in 2023," International Finance Corporation, Washington, DC, Tech. Rep., 2025.